



White Paper

A Tour Beyond BIOS – Memory Protection in UEFI BIOS

*Jiewen Yao
Intel Corporation*

*Vincent J. Zimmer
Intel Corporation*

February 2017

Executive Summary

Introduction

Data execution protection (DEP) is intended to prevent an application or service from executing code from a non-executable memory region. This helps prevent certain exploits that store code via a buffer overflow. [WindowsHeap] shows 4 of 7 exploitation techniques that can be mitigated by DEP and ASLR (Address Space Layout Randomization). [DEP] also shows 14 of 19 exploits from popular exploit kits that fail with DEP enabled. Besides Windows, the Unix/Linux community also has similar non-executable protection [PaX].

In [MemMap], we discussed the DEP and the limitation of enabling the DEP in a UEFI firmware. In [SecurityEnhancement], we only discussed the DEP for stack and setting not-present page for the NULL address and as the guard page. In this document we will have a more comprehensive discussion for the DEP adoption in the current UEFI firmware to harden the pre-boot phase.

Table of Contents

<i>Executive Summary</i>	ii
Introduction.....	ii
<i>Memory Protection in SMM</i>	4
Protection for PE image.....	4
Protection for stack and heap	5
Protection for critical CPU status	5
SMM EntryPoint	5
GDT/IDT	6
Page Table	6
Life cycle of the protection.....	7
SMRAM Size Overhead.....	8
PE image	8
Page Table	8
Performance Overhead	9
Non SMRAM access in SMM	9
Limitation.....	10
Compatibility Consideration.....	10
<i>Memory Protection in UEFI</i>	12
Protection for PE image.....	12
Protection for stack and heap	13
Size Overhead.....	13
Limitation.....	13
Compatibility Consideration.....	14
<i>References</i>	15

Memory Protection in SMM

The SMM is an isolated execution environment [IA32SDM]. [PI] specification volume 4 defines the SMM infrastructure. Figure 1 shows the SMM memory protection. RO means the read-only memory. XD means the execution-disable memory.

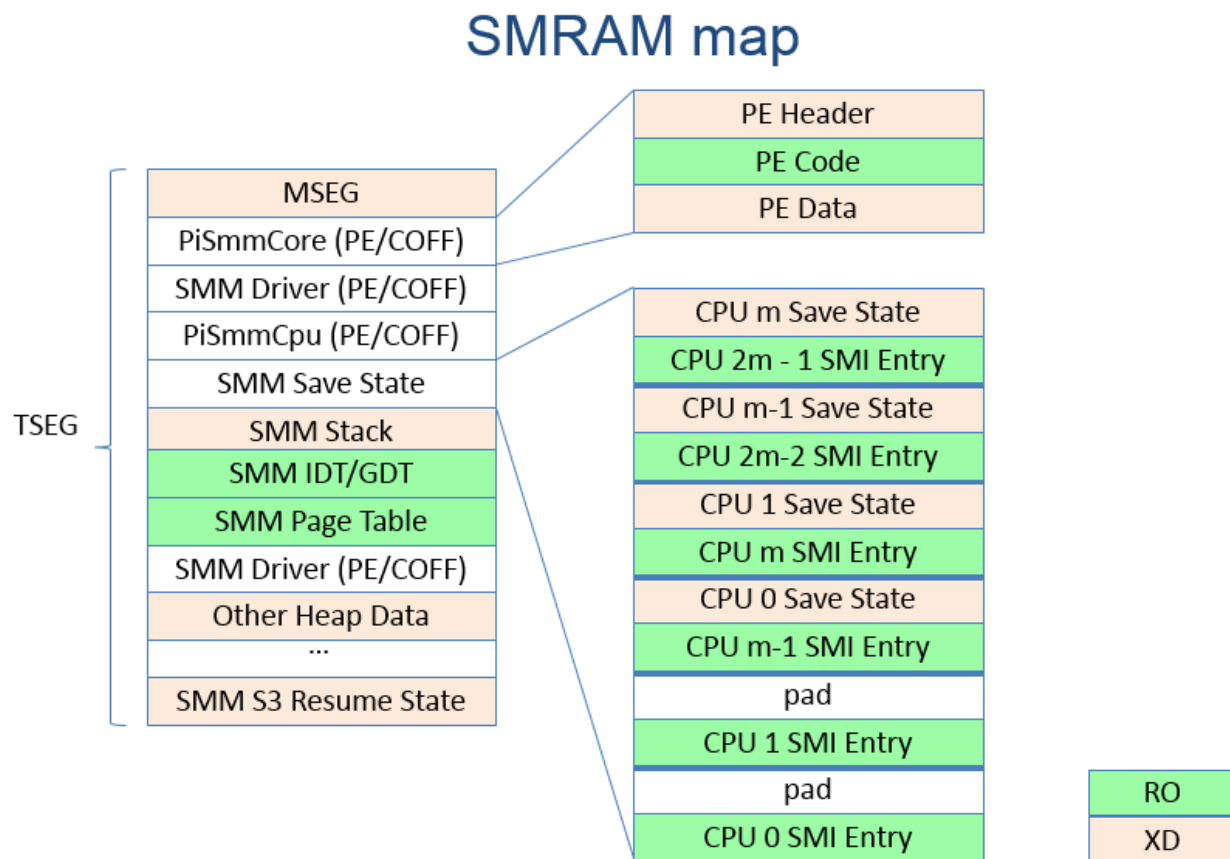


Figure 1 - SMRAM memory protection

Protection for PE image

In UEFI/PI firmware, the SMM image is a normal PE/COFF image loaded by the SmmCore. If the section of SMM image is page aligned, it may be protected according to the section attributes, such as read-only for the code and non-executable for data. See the top right of figure 1.

In EDKII, the PiSmmCore

(<https://github.com/tianocore/edk2/blob/master/MdeModulePkg/Core/PiSmmCore/MemoryAttributesTable.c>) checks the PE image alignment and builds an EDKII_PI_SMM_MEMORY_ATTRIBUTES_TABLE (<https://github.com/tianocore/edk2/blob/master/MdeModulePkg/Include/Guid/PiSmmMemoryAttributesTable.h>) to record such information. If the PI SMM image is not page aligned, this table

will not be published. If `EDKII_PI_SMM_MEMORY_ATTRIBUTES_TABLE` is published, that means the `EfiRuntimeServicesCode` contains the code only and it is `EFI_MEMORY_RO` and the `EfiRuntimeServicesData` contains the data only and it is `EFI_MEMORY_XP`.

Later the `PiSmmCpu` driver

(<https://github.com/tianocore/edk2/blob/master/UefiCpuPkg/PiSmmCpuDxeSmm/SmmCpuMemoryManagement.c>) `SetMemMapAttributes()` API consumes the `EDKII_PI_SMM_MEMORY_ATTRIBUTES_TABLE` and sets the page table attribute.

There are assumptions for the PE image protection in SMM:

- 1) The PE code section and data sections are not merged. If those 2 sections are merged, a `#PF` exception might be generated because the CPU might try to write a RO data item in the data section or execute an NX instruction in code section.
- 2) The PE image can be protected if it is page aligned. There should not be any self-modified-code in the code region. If there is, a platform should not set this PE image to be page aligned.

A platform may disable the XD in the UEFI environment, but this does not impact the SMM environment. The SMM environment may choose to always enable the XD upon SMM entry, and restore the XD state at the SMM exit point.

Protection for stack and heap

The `PiSmmCore` maintains a memory map internally.

(<https://github.com/tianocore/edk2/blob/master/MdeModulePkg/Core/PiSmmCore/Page.c>) If an SMM module allocates the data with `EfiRuntimeServicesCode`, this data is marked as the code page. If the SMM module allocates the data with `EfiRuntimeServicesData`, this data is marked as the data page. This information is also exposed via the `EDKII_PI_SMM_MEMORY_ATTRIBUTES_TABLE`.

The same RO and XP policy is also applied to the normal SMM data region, such as stack and heap.

Protection for critical CPU status

Besides the PE image, the Intel X86 architecture has some special architecture-specific regions that need to be protected as well.

SMM EntryPoint

When a hardware SMI occurs, the Intel X86 CPU jumps to an SMM entry point in order to execute the code at this location. This SMM entry point is not inside of a normal PE image, so we also need to protect this region. See the bottom right of figure 2.

According to [IA32SDM], the SMM entry point is at a fixed offset from `SMBASE`. In EDKII, the `SMBASE` and the SMM save state area are allocated at <https://github.com/tianocore/edk2/blob/master/UefiCpuPkg/PiSmmCpuDxeSmm/PiSmmCpuDxeSmm.c> `PiCpuSmmEntry()`. Both the SMM entry point and the SMM save state are allocated as

a CODE page. Later

<https://github.com/tianocore/edk2/blob/master/UefiCpuPkg/PiSmmCpuDxeSmm/SmmCpuMemoryManagement.c> PatchSmmSaveStateMap() patches the SMM entry point to be read-only and the SMM save state to be non-executable.

GDT/IDT

The GDT defines the base address and the limit of a code segment or a data segment. If the GDT is updated, the code might be redirected to a malicious region. As such, the GDT should be set to read-only.

The IDT defines the entry point of the exception handler. If the IDT is updated, the malicious code may trigger an exception and jump to a malicious region. As such, the IDT should be set to read-only as well.

This work is done by PatchGdtIdtMap() at

<https://github.com/tianocore/edk2/blob/master/UefiCpuPkg/PiSmmCpuDxeSmm/X64/SmmFuncsArch.c>.

However, the IA32 version GDT cannot be set to read-only if the stack guard feature is enabled. (<https://github.com/tianocore/edk2/blob/master/UefiCpuPkg/PiSmmCpuDxeSmm/Ia32/SmmFuncsArch.c>) The reason is that the IA32 stack guard needs to use a “task switch” to switch the stack, and the task switch needs to write the GDT and TSS. The X64 version of the GDT does not have such a problem because the X64 stack guard uses “interrupt stack table (IST)” to switch the stack. For details of the stack switch and exceptions, please refer to [IA32SDM].

Page Table

In an X86 CPU, we rely on the page table to set up the read-only or non-executable region. In order to prevent the page table itself from being updated, we may need to set the page table itself to be read-only.

The work is done at

<https://github.com/tianocore/edk2/blob/master/UefiCpuPkg/PiSmmCpuDxeSmm/X64/PageTbl.c> SetPageTableAttributes().

However, setting a page table to be read-only may break the original dynamic paging feature in SMM. There is a PCD `PcdCpuSmmStaticPageTable` to determine if the platform wants to enable the static page table or the dynamic page table.

If `PcdCpuSmmStaticPageTable` is FALSE, the PiSmmCpu uses the original dynamic paging policy, namely the PiSmmCpu only sets 4GiB paging by default. If the PiSmmCpu needs to access above 4GiB memory locations, a #PF exception is triggered and an above-4GiB mapping is created in the page fault handler.

If `PcdCpuSmmStaticPageTable` is TRUE, the PiSmmCpu will try to set the read-only attribute for the page table.

Figure 2 shows the mapping of the protection.

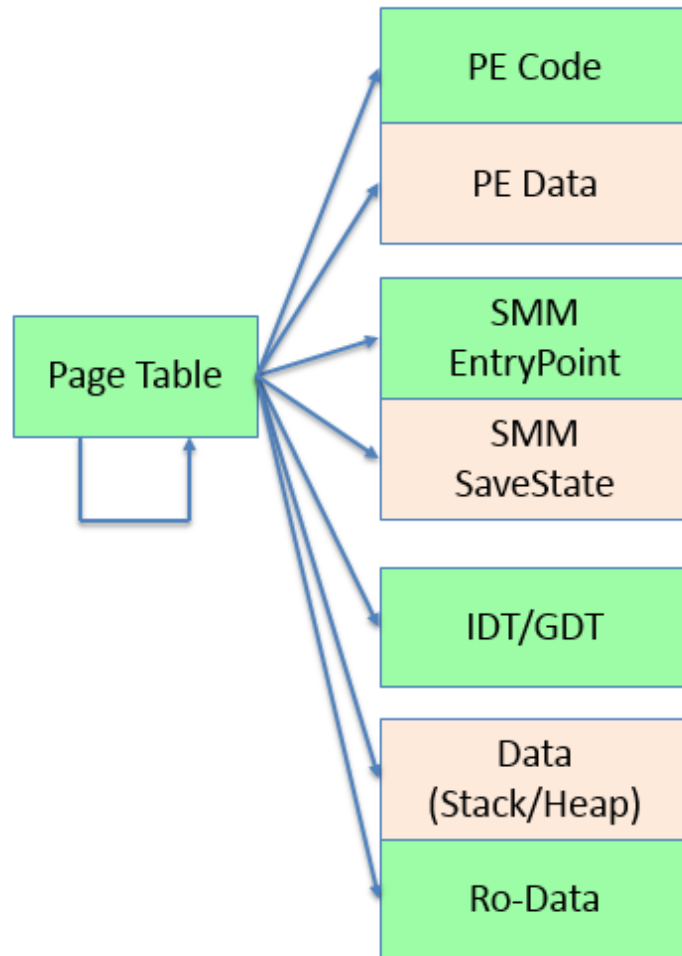


Figure 2 Mapping of Protection in SMM

Life cycle of the protection

In a normal boot, the page table based protection is configured by the PiSmmCpu driver just after the SmmReadyToLock event by PerformRemainingTasks() at

<https://github.com/tianocore/edk2/blob/master/UefiCpuPkg/PiSmmCpuDxeSmm/PiSmmCpuDxeSmm.c>. All read-only data must be ready before SmmReadyToLock.

In an S3 resume, the protection is disabled during SMBASE relocation because the PiSmmCpu needs to set up the environment. The PiSmmCpu uses SmmS3Cr3, which is generated by InitSmmS3Cr3() at

<https://github.com/tianocore/edk2/blob/master/UefiCpuPkg/PiSmmCpuDxeSmm/X64/SmmProfileArch.c> with 4G paging only. After the SMBASE relocation is done, all the protection takes effect up receipt of the next SMI by PerformPreTasks() at <https://github.com/tianocore/edk2/blob/master/UefiCpuPkg/PiSmmCpuDxeSmm/PiSmmCpuDxeSmm.c>.

If there is an additional lock that needs to be set, it can be done in `SmmCpuFeaturesCompleteSmmReadyToLock()` API (defined in <https://github.com/tianocore/edk2/blob/master/UefiCpuPkg/Include/Library/SmmCpuFeaturesLib.h>).

SMRAM Size Overhead

PE image

In order to protect the PE code and data sections, we must set the PE image section alignment to be 4K.

In EDKII, the default PE image alignment is 0x20 bytes. Assuming one PE image has 3 sections (1 header, 1 code section, 1 data section), average overhead for one PE image is $(4K * 3) / 2 = 6K$.

If a platform has n SMM images, the average of the overhead is $6K * n$.

Page Table

In order to protect the page table itself, we must use the static page table instead of the dynamic on-demand page table.

The size of the dynamic paging is fixed. We need 6 fixed pages (24K) and 8 on-demand pages (32K). The total size of the page table is 56K in this case.

The size of the static page table depends upon 2 things: 1) 1G paging capability, 2) max supported address bit. A rough estimation is below:

- 1) If 1G paging is supported,
 - 32 bit addressing need $(1+1+4)$ pages = 24K. (still use 2M paging for below 4G memory)
 - 39 bit addressing need $(1+1+4)$ pages = 24K.
 - 48 bit addressing need $(1+512)$ pages = 2M.
- 2) If 1G paging is not supported, 2M paging is used.
 - 32 bit addressing need $(1+1+4)$ pages = 24K.
 - 39 bit addressing need $(1+1+512)$ pages = 2M.
 - 48 bit addressing need $(1+512+512*512)$ pages = 1G. ⬅ This seems not acceptable.

The maximum address bit is determined by the CPU_HOB if it is present, or the physical address bit returned by the CPUID instruction if the CPU_HOB is not present.

(<https://github.com/tianocore/edk2/blob/master/UefiCpuPkg/PiSmmCpuDxeSmm/X64/PageTbl.c>, `CalculateMaximumSupportAddress()`) A platform may set the CPU_HOB based upon the addressing capability of the memory controller or the CPU.

Performance Overhead

- 1) The SMRAM protection setup is a one-time activity. It happens just after the SmmReadyToLock event. We do not observe too much impact to the system firmware boot performance. The activity only takes some small number of milliseconds.
- 2) The SMRAM runtime protection is based upon the page table. No additional CPU instruction is needed. As such, there is zero SMM runtime performance impact to have this protection.

Non SMRAM access in SMM

Besides the SMRAM, the SMM memory protection also limits the access to the non-SMRAM region.

First, the non-SMRAM region must be set to be non-executable because the SMM entities should not call any code outside SMRAM. Code outside of SMRAM might be controlled by malicious software.

This protection work is done by `InitPaging()` at

<https://github.com/tianocore/edk2/blob/master/UefiCpuPkg/PiSmmCpuDxeSmm/SmmProfile.c>

Second, because of the security concerns regarding SMM entities accessing VMM memory, [WindowsWSMT] and [MicrosoftHV] introduced the Windows SMM Security Mitigations Table (WSMT). A platform needs to report the WSMT table in order to declare that the SMI handler will validate the SMM communication buffer.

As we discussed in [SecureSmmComm], the SMI handler should check if the SMM communication buffer is from a fixed region, (`EfiReservedMemoryType/ EfiACPIMemoryNVS/ EfiRuntimeServicesData/ EfiRuntimeServicesCode`). However, this is a passive check. If a SMI handler does not include such a check, it is hard to detect.

A better way is to use an active check. The `PiSmmCpu` driver sets the non-fixed DRAM region (`EfiLoaderCode/ EfiLoaderData/ EfiBootServicesCode/ EfiBootServicesData/ EfiConventionalMemory/ EfiUnusableMemory/ EfiACPIReclaimMemory`) to be not-present in the page tables after the `SmmReadyToLock` event.

As such, if a platform SMI handler does not include the check recommended in [SecureSmmComm], the system will get `#PF` exception within SMM on such an attack.

This protection work is done by `SetUefiMemMapAttributes()` at

<https://github.com/tianocore/edk2/blob/master/UefiCpuPkg/PiSmmCpuDxeSmm/SmmCpuMemoryManagement.c>.

Figure 3 shows final image layout.

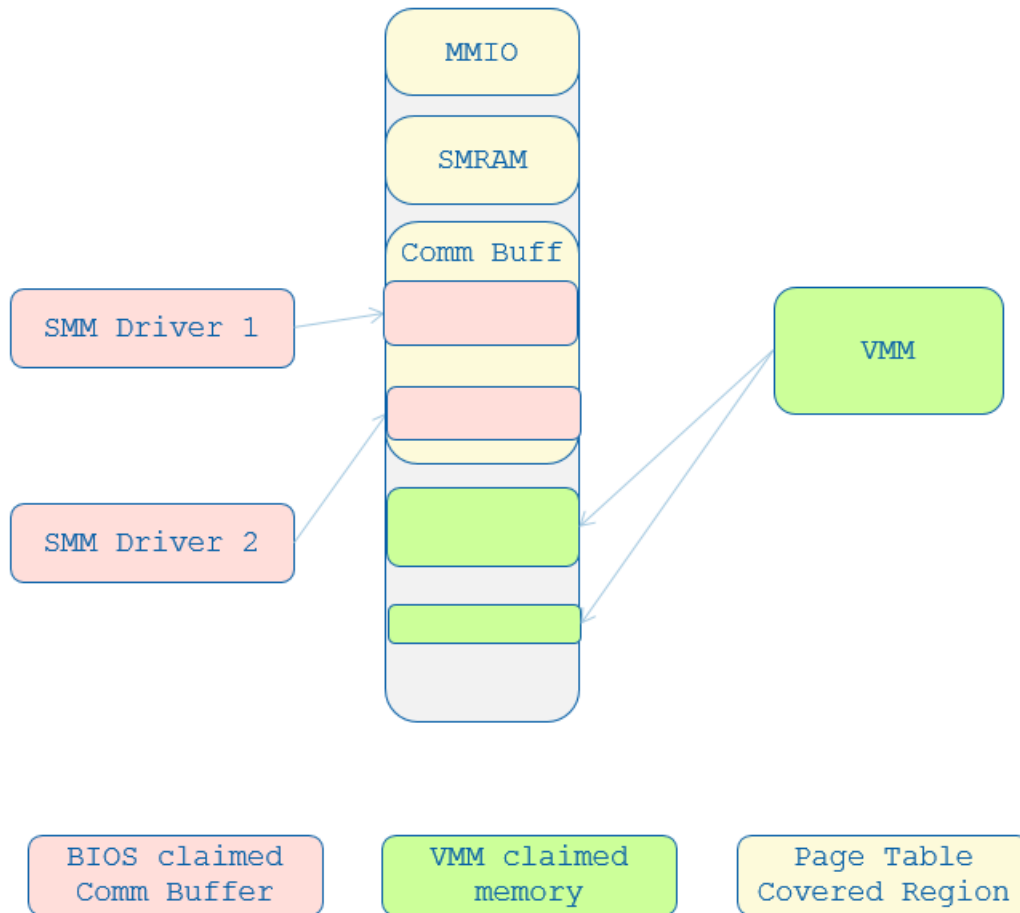


Figure 3 Page table enforced memory layout

Limitation

Setting up RO and NX attribute for SMRAM is a good enhancement to prevent a code overriding attack. However it has some limitations:

- 1) It cannot resist a Return-Oriented-Programming (ROP) attack. [ROP]. We might need ASLR to mitigate the ROP attack. [ASLR] With the code region randomized, an attacker cannot accurately predict the location of instructions in order to leverage gadgets.
- 2) Not all important data structure are set to Read-Only. This is the current SMM driver limitation. The SMM driver can be updated to allocate the important structures to be read-only instead of a read-write global variable.

Compatibility Considerations

- 1) So far, we have not observed self-modified-code in SMM image or executable code in data section. As such, we believe the PE image protection is compatible.
- 2) The protection for the SMM communication buffer may cause a #PF exception in SMM if the SMI handler does not perform the check recommended in [SecureSmmComm].

Summary

This section introduces the memory protection in SMM.

Memory Protection in UEFI

In [MemMap], we discussed to how to report the runtime memory attribute by using `EFI_MEMORY_ATTRIBUTES_TABLE`, so that OS can apply the protection for the runtime code and data. This may bring some compatibility concerns if we choose to adopt the full DEP protection for the entire UEFI memory.

In order to resolve the compatibility concerns, we can define a policy-based setting to enable partial NX and RO protection for the UEFI memory region. The detailed information will be discussed below.

Protection for PE image

The DXE core may apply a pre-defined policy to set up the NX attribute for the PE data region and the RO attribute for the PE code region.

- 1) The image is loaded by the UEFI boot service - `LoadImage()`. If an image is loaded in some other way, the DXE core does not have such knowledge and the DXE core cannot apply any protection.
- 2) The image section is page aligned. If an image is not page aligned, the DXE core cannot apply the page level protection.
- 3) The protection policy can be based upon a PCD `PcdImageProtectionPolicy`. Whenever a new image is loaded, the `DxeCore` checks the source of the image and then decides the policy of the protection. The policy could be to enable the protection if the sections are aligned, or disable the protection. The platform may choose the policy based upon the need. For example, if a platform thinks the image from the firmware volume should be capable of being protection, it can set protection for `IMAGE_FROM_FV`. But if a platform is not sure about a PCI option ROM or a file system on disk, it can set no-protection.

There are assumptions for the PE image protection in UEFI:

- 1) [Same as SMM] The PE code section and data sections are not merged. If those 2 sections are merged, a `#PF` exception might be generated because the CPU may try to write a RO data in data section or execute a NX instruction in the code section.
- 2) [Same as SMM] The PE image can be protected if it is page aligned. There should not be any self-modifying-code in the code region. If there is, a platform should not set this PE image to be page aligned.
- 3) A platform may not disable the XD in the DXE phase. If a platform disables the XD in the DXE phase, the X86 page table will become invalid because the XD bit in page table becomes a RESERVED bit. The consequence is that a `#PF` exception will be generated. If a platform wants to disable the XD bit, it must happen in the PEI phase.

In EDKII, the DXE core checks the image source and alignment, then calls CPU_ARCH->SetMemoryAttribute() for NX or RO whenever a new image is loaded, or an old image is unloaded. When the CPU driver gets the memory attribute setting request, it updates page table.

If an image is loaded before CPU_ARCH protocol is ready, the DXE core just skips the setting, and these images protection will be set in CPU_ARCH callback function.

Protection for stack and heap

[UEFI] specification allows “Stack may be marked as non-executable in identity mapped page tables.” As such, we set up the NX stack

(<https://github.com/tianocore/edk2/blob/master/MdeModulePkg/Core/DxeIplPeim/X64/VirtualMemory.C>, CreateIdentityMappingPageTables()).

Unlike SMM, we cannot set all heap to be NX at this moment because we do not know how the memory is used by the 3rd part vendor. We already observed some unexpected usage in [MemMap], so we do not modify other data attributes.

However, we may use some special techniques, such as the guard page, to apply the protection for the allocated memory in order to detect a buffer overflow. This is discussed in [SecurityEnhancement].

Size Overhead

1) Runtime memory overhead (visible to OS)

The size overhead of the runtime PE image is the same as the overhead of the SMM PE image. If a platform has n runtime images, the average amount overhead is $6K * n$.

2) Boot time memory overhead (invisible to OS)

The size of the overhead for the boot time PE image is the same as the overhead of the SMM PE image.

If a platform has n boot time images, the average overhead is $6K * n$.

The size overhead of the boot time page table is also same as for the SMM static page table. Please refer to the SMM section for the size calculation based upon the 1G paging capability and max supported address bit.

Limitation

The protection in the UEFI is limited to the PE image and the stack at this moment because of the compatibility concerns. The limitations of the UEFI memory protection are:

- 1) Not all images are protected to be NX and RO, based upon the policy.
- 2) Not all heap regions are protected to be NX.

- 3) [Same as SMM] The protection cannot resist ROP attack.
- 4) [Same as SMM] Not all important data structures are set to ReadOnly.

Compatibility Consideration

A platform may need to evaluate and select the image protection policy based upon the capability of the platform image, Option ROM, and OS loader. For platform images, the Compatibility Support Module (CSM) and the EDK-I Compatibility Package (ECP) modules should be considered. If a platform observes the compatibility issues, it should choose 1) to disable the protection, or 2) to fix the compatibility issue and enable the protection.

Summary

This section introduces the memory protection in UEFI.

References

[ASLR] *Address Space Layout Randomization*,
https://en.wikipedia.org/wiki/Address_space_layout_randomization

[DEP] *Exploit Mitigation Improvements in Windows 8*, Ken Johnson, Ma, Miller,
http://media.blackhat.com/bh-us-12/Briefings/M_Miller/BH_US_12_Miller_Exploit_Mitigation_Slides.pdf

[IA32SDM] *Intel® 64 and IA-32 Architectures Software Developer's Manual*, www.intel.com

[MemMap] *A Tour Beyond BIOS Memory Map And Practices in UEFI BIOS*, Jiewen Yao, Vincent Zimmer, 2016 https://github.com/tianocore-docs/Docs/raw/master/White_Papers/A_Tour_Beyond_BIOS_Memory_Map_And_Practices_in_UEFI_BIOS_V2.pdf

[PaX] *PaX Home Page*, <https://pax.grsecurity.net/>

[ROP] *Return-oriented programming*, https://en.wikipedia.org/wiki/Return-oriented_programming

[SecureSmmComm] *A Tour Beyond BIOS Secure SMM Communication*, Jiewen Yao, Vincent Zimmer, Star Zeng, 2016, https://github.com/tianocore-docs/Docs/raw/master/White_Papers/A_Tour_Beyond_BIOS_Secure_SMM_Communication.pdf

[SecurityEnhancement] *A Tour Beyond BIOS Security Enhancement to Mitigate Buffer Overflow in UEFI*, Jiewen Yao, Vincent Zimmer, 2016, https://github.com/tianocore-docs/Docs/raw/master/White_Papers/A_Tour_Beyond_BIOS_Security_Enhancement_to_Mitigate_Buffer_Overflow_in_UEFI.pdf

[SecurityDesign] *A Tour Beyond BIOS Security Design Guide in EDKII*, Jiewen Yao, Vincent Zimmer, 2016, https://github.com/tianocore-docs/Docs/raw/master/White_Papers/A_Tour_Beyond_BIOS_Security_Design_Guide_in_EDKII.pdf

[UEFI] *Unified Extensible Firmware Interface (UEFI) Specification, Version 2.6*
www.uefi.org

[WindowsHeap] *Preventing the exploitation of user mode heap corruption vulnerabilities*, 2009,
<https://blogs.technet.microsoft.com/srd/2009/08/04/preventing-the-exploitation-of-user-mode-heap-corruption-vulnerabilities/>

[WindowsInternal] *Windows Internals, 6th edition*, Mark E. Russinovich, David A. Solomon, Alex Ionescu, 2012, Microsoft Press. ISBN-13: 978-0735648739/978-0735665873

[WindowsWSMT] *Windows SMM Security Table*, [https://msdn.microsoft.com/en-us/library/windows/hardware/dn495660\(v=vs.85\).aspx#wsmt](https://msdn.microsoft.com/en-us/library/windows/hardware/dn495660(v=vs.85).aspx#wsmt)
<http://download.microsoft.com/download/1/8/A/18A21244-EB67-4538-BAA2-1A54E0E490B6/WSMT.docx>

[MicrosoftHV] *Microsoft Hypervisor Requirements*, <https://msdn.microsoft.com/en-us/library/windows/hardware/dn614617>

Authors

Jiewen Yao (jiewen.yao@intel.com) is EDKII BIOS architect, EDKII FSP package maintainer, EDKII TPM2 module maintainer, EDKII ACPI S3 module maintainer, with Software and Services Group at Intel Corporation. Jiewen is member of UEFI Security Sub-team and PI Security Sub-team in the UEFI Forum.

Vincent J. Zimmer (vincent.zimmer@intel.com) is a Senior Principal Engineer with the Software and Services Group at Intel Corporation. Vincent chairs the UEFI Security and Network Sub-teams in the UEFI Forum.

This paper is for informational purposes only. THIS DOCUMENT IS PROVIDED "AS IS" WITH NO WARRANTIES WHATSOEVER, INCLUDING ANY WARRANTY OF MERCHANTABILITY, NONINFRINGEMENT, FITNESS FOR ANY PARTICULAR PURPOSE, OR ANY WARRANTY OTHERWISE ARISING OUT OF ANY PROPOSAL, SPECIFICATION OR SAMPLE. Intel disclaims all liability, including liability for infringement of any proprietary rights, relating to use of information in this specification. No license, express or implied, by estoppel or otherwise, to any intellectual property rights is granted herein.

Intel, the Intel logo, Intel. leap ahead. and Intel. Leap ahead. logo, and other Intel product name are trademarks or registered trademarks of Intel Corporation or its subsidiaries in the United States and other countries.

*Other names and brands may be claimed as the property of others.

Copyright 2017 by Intel Corporation. All rights reserved

